

# Predicting water potability using machine learning

Midterm Project Presentation

WOO-CHAN KIM

# **I. Project title**

## **Predicting Water Potability Using Machine Learning**

# **II. Project introduction**

## **1) Objective**

: Developing a machine learning model to predict water potability based on various water quality indicators.

## **2) Motivation**

- (1) I chose this project because it is meaningful work that would make a social contribution like public health improvement, environmental protection and sustainability .
- (2) It has correlations with my first project of movie review classifier. So, I could get a deeper knowledge of classification.

# III. Analysis of original code

## 1) Dataset

- Train: 1,608
- Validation: 201
- Test: 202

## 2) Machine Learning Model

: Xgboost

## 3) Performance

- Accuracy : 0.6584
- F1-score(0): 0.74
- F1-score(1): 0.51

```
y_pred = best_xgb.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Test Accuracy: {accuracy:.4f}")
```



Test Accuracy: 0.6584

```
# Generate and print the classification report
report = classification_report(y_test, y_pred)
print(f"Classification Report:\n{report}")
```

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.68      | 0.81   | 0.74     | 120     |
| 1            | 0.61      | 0.44   | 0.51     | 82      |
| accuracy     |           |        | 0.66     | 202     |
| macro avg    | 0.64      | 0.62   | 0.62     | 202     |
| weighted avg | 0.65      | 0.66   | 0.65     | 202     |

Origin Code



## IV. Updating Preprocessing Process

### 1) Dealing with missing values

: Replace missing values with a representative value rather than deleting them.

#### III. Dataset Description

##### Water Potability Dataset

In this project, a total of 3,276 data samples are preprocessed and 2011 data are divided into training sets, validation sets, and test sets for model development.

##### Training Set:

Composed of 1,608 samples. This dataset is used to train the machine learning model. The model learns the relationship between various water quality indicators and potability during this stage.

##### Validation Set:

Composed of 201 samples. This set is used to evaluate model performance during training, allowing for tuning of hyperparameters and preventing overfitting. It helps monitor the model's generalization performance and adjust the learning process accordingly.

##### Test Set:

Composed of 202 samples. After the model is fully trained, this set is used to assess the final performance of the model. It provides an evaluation of how well the model predicts water potability in real-world scenarios.

- **Origin Dataset:** 3,276
- **Sum of train/validation/test sets:** 2,011
- **The number of missing data:** 1,265

```
# Check the missing value percentage
missing_values = df.isnull().sum()
missing_percentage = (df.isnull().sum() / len(df)) * 100

# Generate Missing value info table
missing_summary = pd.DataFrame({
    'Missing Values': missing_values,
    'Percentage (%)': missing_percentage
})

print(missing_summary)
```

|                 | Missing Values | Percentage (%) |
|-----------------|----------------|----------------|
| ph              | 491            | 14.987790      |
| Hardness        | 0              | 0.000000       |
| Solids          | 0              | 0.000000       |
| Chloramines     | 0              | 0.000000       |
| Sulfate         | 781            | 23.840049      |
| Conductivity    | 0              | 0.000000       |
| Organic_carbon  | 0              | 0.000000       |
| Trihalomethanes | 162            | 4.945055       |
| Turbidity       | 0              | 0.000000       |
| Potability      | 0              | 0.000000       |

```
[7] # Replace missing values by Mean
df['ph'] = df['ph'].fillna(df['ph'].mean())

[8] # Replace missing values by Mean
df['Sulfate'] = df['Sulfate'].fillna(df['Sulfate'].mean())

[9] # Replace missing values by Mean
df['Trihalomethanes'] = df['Trihalomethanes'].fillna(df['Trihalomethanes'].mean())
```

Replace missing value  
to mean value

Detecting missing values

# IV. Updating Preprocessing Process

## 2) Dealing with outlier values

```
[ ] from scipy import stats





outlier_indices = set()

# Z-Score
for column in df.select_dtypes(include='number').columns:
    z_scores = np.abs(stats.zscore(df[column]))
    outliers = df[z_scores > 3]
    outlier_indices.update(outliers.index)
    print(outliers)

print("Outlier Index:", outlier_indices)
print(f"The number of Outliers: {len(outlier_indices)}")

df = df.drop(index=outlier_indices)
```

Code



Before deleting outliers

After deleting outliers

# IV. Updating Preprocessing Process

## 3) Data Normalization & Splitting the dataset

```
[47] from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
normalized_data = scaler.fit_transform(df)

df = pd.DataFrame(normalized_data, columns=df.columns)

print("Normalized Dataset: ")
print(df)
```

```
Normalized Dataset:
   ph  Hardness  Solids  Chloramines  Sulfate  Conductivity \
0  0.501200  0.561805  0.427759  0.524547  0.661087  0.779857
1  0.117052  0.163786  0.382596  0.453558  0.500481  0.841303
2  0.617462  0.663836  0.409333  0.735459  0.500481  0.466566
3  0.642310  0.611819  0.453401  0.605586  0.607320  0.347575
4  0.730844  0.436341  0.368991  0.444095  0.391193  0.423142
...
3123 0.225744  0.502690  0.987569  0.510287  0.621478  0.698398
3124 0.584322  0.502012  0.355426  0.605803  0.500481  0.410324
3125 0.768210  0.408184  0.686129  0.529887  0.500481  0.495462
3126 0.278109  0.697418  0.243714  0.418127  0.500481  0.432758
3127 0.591836  0.510182  0.356980  0.546868  0.500481  0.270582

   Organic_carbon  Trihalomethanes  Turbidity  Potability
0  0.306829  0.722415  0.288245  0.0
1  0.551982  0.395786  0.623577  0.0
2  0.638222  0.503282  0.308484  0.0
3  0.718296  0.864635  0.651519  0.0
4  0.367016  0.136596  0.530758  0.0
...
3123 0.486325  0.506132  0.609436  1.0
3124 0.793202  0.503028  0.252282  1.0
3125 0.340500  0.539770  0.361469  1.0
3126 0.347133  0.621186  0.668942  1.0
3127 0.601028  0.634078  0.145611  1.0
```

Data Normalization Code

```
[63] from sklearn.model_selection import train_test_split

# Train:Validation:Test = 75:10:15
X_train, X_temp, y_train, y_temp = train_test_split(
    X_scaled, y, test_size=0.25, random_state=42, stratify=y)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.6, random_state=42, stratify=y_temp)

print("Train set size:", X_train.shape, y_train.shape)
print("Validation set size:", X_val.shape, y_val.shape)
print("Test set size:", X_test.shape, y_test.shape)
print("\nTrain set class distribution:\n", y_train.value_counts())
print("\nValidation set class distribution:\n", y_val.value_counts())
print("\nTest set class distribution:\n", y_test.value_counts())
```

```
Train set size: (2346, 9) (2346,)
Validation set size: (312, 9) (312,)
Test set size: (470, 9) (470,)

Train set class distribution:
Potability
0    1447
1     899
Name: count, dtype: int64

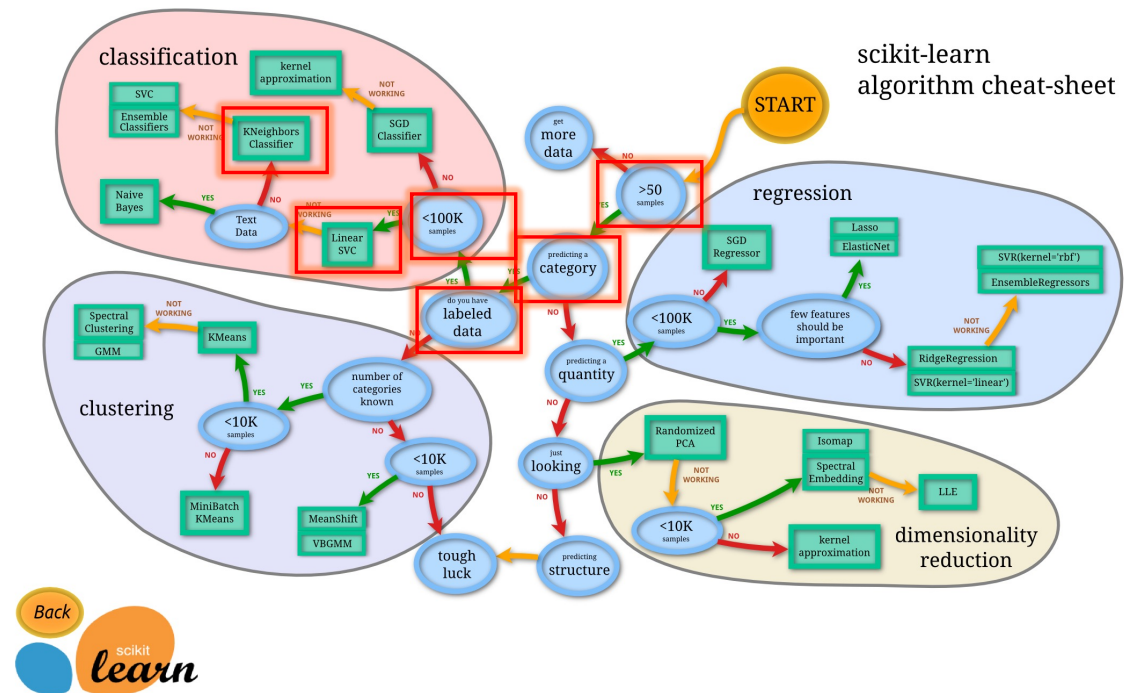
Validation set class distribution:
Potability
0     193
1     119
Name: count, dtype: int64
```

Data Splitting Code

## V. Updating the model

### The Most Suitable Models for this Project

- Logistic Regression
- SVC / SVM
- Randomforest



# V. Updating the model

## 1) Logistic Regression

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Initialization
lr_model = LogisticRegression(max_iter=1000)

# Model training
lr_model.fit(X_train, y_train)
```

Training code

```
# Validation dataset prediction
y_val_pred = lr_model.predict(X_val)
val_accuracy = accuracy_score(y_val, y_val_pred)
print(f'Validation Accuracy: {val_accuracy:.4f}')
print('Classification Report:')
print(classification_report(y_val, y_val_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_val, y_val_pred))
```

Validation Accuracy: 0.6186

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.62      | 1.00   | 0.76     | 193     |
| 1            | 0.00      | 0.00   | 0.00     | 119     |
| accuracy     |           |        | 0.62     | 312     |
| macro avg    | 0.31      | 0.50   | 0.38     | 312     |
| weighted avg | 0.38      | 0.62   | 0.47     | 312     |

Confusion Matrix:

```
[[193  0]
 [119  0]]
```

**Validation set score: 0.6186**



# V. Updating the model

## 2) SVC / SVM

```
from sklearn.svm import SVC

svm_model = SVC(kernel='rbf', probability=True, random_state=42)

svm_model.fit(X_train, y_train)

y_val_pred = svm_model.predict(X_val)
val_accuracy = accuracy_score(y_val, y_val_pred)
print(f'Validation Accuracy: {val_accuracy:.4f}')
print('Classification Report:')
print(classification_report(y_val, y_val_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_val, y_val_pred))
```

```
Validation Accuracy: 0.6731
Classification Report:
      precision    recall  f1-score   support

     0       0.67      0.93      0.78       193
     1       0.70      0.25      0.37       119

 accuracy      0.68
 macro avg     0.68
 weighted avg   0.67

Confusion Matrix:
[[180  13]
 [ 89  30]]
```

Training & Validation score code

```
Test Accuracy: 0.6681
Classification Report:
      precision    recall  f1-score   support

     0       0.71      0.79      0.75       290
     1       0.58      0.48      0.52       180

 accuracy      0.67
 macro avg     0.64
 weighted avg   0.66

Confusion Matrix:
[[228  62]
 [ 94  86]]
```

**Test set score: 0.6681**

# V. Updating the model

## 3) Randomforest

```
from sklearn.ensemble import RandomForestClassifier

rf_model = RandomForestClassifier(n_estimators=100, random_state=42, class_weight='balanced')
rf_model.fit(X_train, y_train)

y_val_pred = rf_model.predict(X_val)
val_accuracy = accuracy_score(y_val, y_val_pred)
print(f'Validation Accuracy: {val_accuracy:.4f}')
print('Classification Report:')
print(classification_report(y_val, y_val_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_val, y_val_pred))
```

Validation Accuracy: 0.6667  
Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.67      | 0.91   | 0.77     | 193     |
| 1            | 0.65      | 0.27   | 0.38     | 119     |
| accuracy     |           |        | 0.67     | 312     |
| macro avg    | 0.66      | 0.59   | 0.58     | 312     |
| weighted avg | 0.66      | 0.67   | 0.62     | 312     |

Confusion Matrix:  
[[176 17]  
[ 87 32]]

Training & Validation score code

```
y_test_pred = rf_model.predict(X_test)
test_accuracy = accuracy_score(y_test, y_test_pred)
print(f'Test Accuracy: {test_accuracy:.4f}')
```

Classification Report:

```
print(classification_report(y_test, y_test_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_test_pred))
```

Test Accuracy: 0.6617  
Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.66      | 0.92   | 0.77     | 290     |
| 1            | 0.65      | 0.25   | 0.36     | 180     |
| accuracy     |           |        | 0.66     | 470     |
| macro avg    | 0.66      | 0.58   | 0.57     | 470     |
| weighted avg | 0.66      | 0.66   | 0.61     | 470     |

Confusion Matrix:  
[[266 24]  
[135 45]]

Test set score: 0.6617

## VI. Result

### [ Original Code ]

#### 1) Dataset

- Train: 1,608
- Validation: 201
- Test: 202

#### 2) Machine Learning Model

- Xgboost

#### 3) Performance

- Accuracy : 0.6584



### [ Updated Code ]

#### 1) Dataset

- Train: 2,346
- Validation: 312
- Test: 470

#### 2) Machine Learning Model

- Logistic Regression
- SVC / SVM
- Randomforest

#### 3) Performance

- Highest Accuracy : 0.6681

# Q & A

E-mail : [kimwc620@korea.ac.kr](mailto:kimwc620@korea.ac.kr)

GitHub address :

<https://github.com/SkyDreamer14/WaterPotability>